

Information Retrieval + Q&A

(SNLP Tutorial 12)

Vilém Zouhar, Awantee Deshpande, Julius Steuer

13th July, 15th July

Evaluation metrics

- Documents D , queries Q
- System: $Q \rightarrow \mathcal{P}(D)$
- For $q \in Q$: retrieved (output), relevant (gold)
- Recall $\frac{|\text{retrieved} \cap \text{relevant}|}{|\text{relevant}|}$
- Precision $\frac{|\text{retrieved} \cap \text{relevant}|}{|\text{retrieved}|}$
- $\{\text{Precision}, \text{Recall}\}@k$: Retrieve k documents (top k scoring)
- Recall@ k $\frac{|\text{retrieved}@k \cap \text{relevant}|}{|\text{relevant}|}$
- Precision@ k $\frac{|\text{retrieved}@k \cap \text{relevant}|}{k}$

Evaluation metrics

- Average precision: $AveP(q) = \frac{\sum_k P@k \times rel(k)}{|relevant|}$
- $rel(k) = \begin{cases} 1 & k\text{-th document relevant} \\ 0 & otherwise \end{cases}$
- Mean average precision $MAP(Q) = \frac{\sum_{q \in Q} AveP(q)}{|Q|}$
 - ▶ Q can be a “test set”
- F-score $2 \cdot \frac{P \cdot R}{P + R}$
- F-score@k $2 \cdot \frac{P@k \cdot R@k}{P@k + R@k}$

Evaluation metrics

Taking the rank into consideration

- Mean Reciprocal Rank
- $MRR(Q) = \frac{1}{|Q|} \sum_{q \in Q} \frac{1}{\text{rank}_q}$
rank_q = position of the first relevant document

document	rank	relevant
a	4	+
b	1	
c		
d		+
e	2	+
f	3	

- $Q = \{\text{example}\}, MRR(Q) = \frac{1}{\text{rank}_{\text{example}}} = \frac{1}{2}$

Document Retrieval - example

- Query: **Goethe, devil**
- Document:
 - A: Wolfgang's idea of the demon Mephistopheles who makes a bet with God
 - B: Faust is Wolfgang **Goethe**'s play in German about a pact with the **devil**
 - C: **Devil**ishly good lasagne
 - D: The impact of **Goethe**'s demon play on the German literature
- How to rank them?
 - B (contains the two key words)
 - D (Goethe, literature)
 - A (Wolfgang - Goethe, Mephistopheles - devil)
 - C (unrelated context)
- Can these inferences be made automatically?

Document Retrieval - Bag of Words

- Text must be represented as a vector of numbers
- BoW model requires: i) Vocabulary, ii) Measure of presence of words
- e.g. Vocabulary = {'to', 'be', 'or', 'not', 'question'}

Document: *to be or not to be*

BoW representation: {to:2, be:2, or:1, not:1} \rightarrow [1 1 1 1 0]

- Can also store counts
- Disregard grammar, word order

Solution 1 (counts)

- Solution: vector with counts of words:
(<the>, <a>, <dog>, <president>, ...)
(57, 68, 0, 2, ...)
- Issue: representation vectors are enormous
- Issue: longer documents have naturally higher counts
- Issue: useless stop words

Solution 2 (tf)

- Solution: vector with counts of non-stop words, normalized by total words:
(`<dog>`, `<president>`, `<princess>`, `<thing>`, ...)
(0, 0.0003, 0.00001, 0.08, ...).
- Issue: some words naturally occur with higher frequency but don't contribute to document meaning (`<thing>`)
- Issue: how do we know which words are useful?

Term Frequency - Inverse Document Frequency

TF-IDF

$$tf(term, doc) = \frac{count_{doc}(term)}{|doc|}$$

$$idf'(term) = \frac{|D|}{df(term)}, idf(term) = \log_2 \left(\frac{|D|}{df(term)} \right)$$

$$tf-idf(term, doc) = tf(term, doc) \times idf(term)$$

Augmented TF

$$tf'(term, doc) = 0.5 + 0.5 \cdot \frac{count_{doc}(term)}{\max_{term'} \{count_{doc}(term')\}}$$

└ Term Frequency - Inverse Document Frequency

TF-IDF

$$tf(term, doc) = \frac{count_{doc}(term)}{|doc|}$$

$$idf(term) = \frac{|D|}{df(term)}, idf(term) = \log_2 \left(\frac{|D|}{df(term)} \right)$$

$$tf \cdot idf(term, doc) = tf(term, doc) \times idf(term)$$

Augmented TF

$$tf'(term, doc) = 0.5 + 0.5 \cdot \frac{count_{doc}(term)}{\max_{term'} \{count_{doc}(term')\}}$$

- Probability that i-th term occurs k times in the document: $p_{\lambda_i}(k) = e^{-\lambda_i} \frac{\lambda_i^k}{k!}$ (λ_i parameter of the distribution)
- Expected value of occurrence: $N \cdot E_i(k) = N \cdot \lambda_i = \text{collection frequency}_i$
- Term present at least once: $N \cdot (1 - P_{\lambda_i}(0)) = \text{document frequency}_i$

Solution 3

- Solution: vector of tf-idf
- Good metrics to determine the significance of a term in a document collection
- Issue: still enormous vectors
- Issue: demon – Mephistopheles are equally separate concepts as demon – lasagne
- Issue: independent terms assumption

Document Retrieval - Probabilistic Retrieval

- Goal: Find $P[R|d,q]$
- Ranking: Proportional to relevance odds

$$O(R|d) = \frac{P[R|d]}{P[\bar{R}|d]}$$

- Estimation results in tf-idf with logarithmically damped idfs

Document retrieval - Language Model

- Pretend the query was generated by a LM based on the document
 - Ranking: Proportional to query likelihood
 - $\operatorname{argmax}_d p(d|q) = \operatorname{argmax}_d \frac{p(q|d) \cdot p(d)}{p(q)} = \operatorname{argmax}_d p(q|d) \cdot p(d)$
 - $\approx \operatorname{argmax}_d p_{LM}(q|d) \cdot p(d)$
 - $p(d) \approx \frac{1}{|D|}$
 - $\approx \operatorname{argmax}_d p_{LM}(q|d)$
 - Unigram: $p(d|q) \approx \prod_i p_{LM}(q_i|d)$
 - Jelinek-Mercer smoothing [9]: $p(q_i|d, D) = \lambda \cdot p(q_i|d) + (1 - \lambda) \cdot p(q_i|D)$
 - High λ : documents with all query words (conjunctive)
 - Low λ : suitable for long queries (disjunctive)
 - Issue: Without word embeddings, no word relatedness
- Query: **Goethe, devil**
- A: Wolfgang's idea of the demon Mephistopheles who makes a bet with God

└ Document retrieval - Language Model

- Pretend the query was generated by a LM based on the document
- Ranking: Proportional to query likelihood
- $\text{argmax}_d p(d|q) = \text{argmax}_d \frac{p(q|d)p(d)}{p(q)} \approx \text{argmax}_d p(q|d) \cdot p(d)$
- $\approx \text{argmax}_d p_{LM}(q|d) \cdot p(d)$
- $p(d) \approx \frac{1}{|D|}$
- $\approx \text{argmax}_d p_{LM}(q|d)$
- Unigram: $p(d|q) \approx \prod_i p_{LM}(q_i|d)$
- Jelinek-Mercer smoothing [0]: $p(q_i|d, D) = \lambda \cdot p(q_i|d) + (1 - \lambda) \cdot p(q_i|D)$
- High λ : documents with all query words (conjunctive)
- Low λ : suitable for long queries (disjunctive)
- Issue: Without word embeddings, no word relatedness
- Query: **Goethe, devil**
- A: Wolfgang's idea of the demon Mephistopheles who makes a bet with God

- Other smoothing schemas exist, like discounting, adding epsilon or linear interpolation between multiple LMs, including zerogram
- Other improvements, such as special grammar, prior knowledge of the document (length), list of synonyms, etc

Document vector representation

- Represent the query and all documents as a vector
Measure their similarity (L-norm, cosine distance: $\frac{D \cdot Q}{|D||Q|}$)
- How to represent a query/document as a fixed size vector? Can we model word co-occurrence for a topic?

Solution 4 (LSA)

- Assumption: Documents are composed of k latent topics
- Solution: Perform dimensionality reduction using SVD
- \rightarrow eigenvalues, singular value decomposition
- $A_{i,j} = \#$ occurrences of term t_i in document d_j (replace with tf-idf later)

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	0	0	0
Faust	0	1	0	0
Goethe	0	1	0	1
devil	0	1	1	0
demon	1	0	0	1
lasagne	0	0	1	0
German	0	1	0	1

└ Solution 4 (LSA)

Solution 4 (LSA)

- Assumption: Documents are composed of k latent topics
- Solution: Perform dimensionality reduction using SVD
- \rightarrow eigenvalues, singular value decomposition
- $A_{i,j} = \#$ occurrences of term t_i id document d_j (replace with tf-idf later)

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	0	0	0
Faust	0	1	0	0
Gotha	0	1	0	1
devil	0	1	1	0
demon	1	0	0	1
Isagane	0	0	1	0
German	0	1	0	1

The example uses counts, but for better representation of term importance in the document, one would use tf-idf.

Approximation of A

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	1	0	1
Faust	1	1	0	0
Goethe	1	1	0	0
devil	1	1	0	1
demon	1	1	0	1
lasagne	0	0	1	0
German	1	1	0	0

	c_1	c_2	c_3
Wolfgang	1	0	0
Mephistopheles	0	1	0
Faust	1	0	0
Goethe	1	0	0
devil	0	1	0
demon	0	1	0
lasagne	0	0	1
German	1	0	0

3 latent concepts:

{Goethe (Wolfgang, Faust, German), devil (Mephistopheles, demon), lasagne}

$$d_1 = 1 \times c_1 + 1 \times c_2$$

Information Retrieval + Q&A

└ Approximation of A

TODO

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	1	0	1
Faust	1	1	0	0
Goethe	1	1	0	0
devil	1	1	0	1
demon	1	1	0	1
lasagne	0	0	1	0
German	1	1	0	0

	c_1	c_2	c_3
Wolfgang	1	0	0
Mephistopheles	0	1	0
Faust	1	0	0
Goethe	1	0	0
devil	0	1	0
demon	0	1	0
lasagne	0	0	1
German	1	0	0

3 latent concepts
 {Goethe (Wolfgang, Faust, German), devil (Mephistopheles, demon), lasagne}

$$d_1 = 1 \times c_1 + 1 \times c_2$$

Approximation of A

- Given: A, k
- $A' = \operatorname{argmin}_{A' \text{ rank } k} \|A - A'\|$
- Distance e.g. Frobenius ($\sqrt{\sum_{i,j} a_{i,j}^2}$)

└ Approximation of A

Given k concepts, we may wish to find such a matrix A' , that's as close to the original one, but with every document being a combination of k independent vectors.

- Given: A, k
- $A' = \underset{A' \in \mathbb{R}^{n \times k}}{\operatorname{argmin}} \|A - A'\|$
- Distance e.g. Frobenius ($\sqrt{\sum_{i,j} A_{ij}^2}$)

SVD

- $A_{i,j} = \#$ occurrences of term t_i in document d_j (replace with tf-idf later)
- $(A^T A)_{i,j} = \#$ intersection of documents d_i and d_j
- $(A A^T)_{i,j} = \#$ documents in which both terms t_i and t_j occur (multiplied counts)
- U = eigenvectors of $A^T A$
- V = eigenvectors of $A A^T$
- S = roots of corresponding eigenvalues of $A^T A$
- $A = U S V^T$

Eigen{vector,value}

Nonzero $v \in \mathbb{R}^n, \lambda \in \mathbb{R}$

Eigenvector

$$Av = \lambda v \quad Av = \lambda Iv \quad (A - \lambda I)v = 0 \quad \ker(A - \lambda I)$$

"Directions (v) which A only scales."

Eigenvalue

$$Av = \lambda v$$

"The stretch (λ) of eigenvector v by A ."

Proof sketch

$$A = USV^T, A^T = VSU^T, S \text{ diagonal}$$

$$U^T U = VV^T = I \text{ orthogonal}$$

$$AA^T U = US^2 \rightarrow U \text{ eigenvectors of } AA^T, S \text{ root of eigenvalues}$$

$$(\forall i : AA^T U_{i,*} = U_{i,*} \cdot S_{i,i}^2)$$

$$A^T AV = VS^2 \rightarrow V \text{ eigenvectors of } A^T A, S \text{ root of eigenvalues}$$

$$(\forall i : A^T AV_{i,*} = V_{i,*} \cdot S_{i,i}^2)$$

- ① Order eigenvalues by descending values ($S_{i,i} > S_{i+1,i+1} \geq 0$)
(proof next slide)
- ② Take top-k eigenvectors + values (or all above threshold)
- ③ $A_K = U_K S_K V_K^T$ $[(m \times n), (n \times n), (n \times n)] \rightarrow [(m \times k), (k \times k), (k \times n)]$
 - Term \rightarrow latent representation: $U_k S_k$
 - Document \rightarrow latent representation: $(S_k V_k^T)^T = V_k S_k^T = V_k S_k$

└ LSA

- Order eigenvalues by descending values ($S_{i,j} > S_{i+1,j+1} \geq 0$)
(proof next slide)
- Take top-k eigenvectors + values (or all above threshold)
- $A_k = U_k S_k V_k^T$ $[(m \times n), (n \times n)] \rightarrow [(m \times k), (k \times k), (k \times n)]$
 - Term \rightarrow latent representation: $U_k S_k$
 - Document \rightarrow latent representation: $(S_k V_k^T)^T = V_k S_k^T = V_k S_k$

- We are free to permute the eigenvalues, so we can order them (together with the vectors) and also we know that the eigenvalues are non-negative
- Therefore we can just take the top-k eigenvalues and replace the rest with zero.
- Essentially this crops the neighbouring matrices to first k columns and first k rows of V^T .

Properties of S

Descending

$U' = U$ +swapped i, j column, $S' = S$ +swapped i, j values, $V'^T = V^T$ +swapped i, j row

$$U' = U \times C(i, j), S' = S \times C(i, j), V'^T = V^T \times R(i, j)$$

$$U'S' = (US) \text{ with swapped } i, j \text{ columns, } U'S' = (US) \times C(i, j)$$

$$U'S'V'^T = (US) \times C(i, j) \times V^T \times R(i, j) = (US) \times C(i, j) \times C(i, j)V^T = USV^T$$

Non-negative

$A^T A$ is positive semidefinite $\Rightarrow S_{i,i} \geq 0$

$$\forall x \neq \vec{0} : x^T A^T A x = (Ax)^T (Ax) = \|Ax\|^2 \geq 0$$

LSA Concepts

- $U_k S_k$ maps terms to latent “concepts” ($m \rightarrow k$)
- $V_k S_k$ maps documents to “concepts” ($n \rightarrow k$)

└ LSA Concepts

- $U_k S_k$ maps terms to latent "concepts" ($m \rightarrow k$)
- $V_k S_k$ maps documents to "concepts" ($n \rightarrow k$)

- The k then becomes obvious is the number of concepts
- We don't specify the concepts, they are determined by SVD
- From our point of view, they are latent

LSA Example

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	0	0	0
...				

- Choose $k = 2$
- Representation of Goethe: fourth row of U_k ($m \times k \rightarrow 1 \times 2$) scaled by S_k : $[0.13, -0.13]^T$
- Representation of devil: fifth row of U_k ($m \times k \rightarrow 1 \times 2$) scaled by S_k : $[0.58, -0.01]^T$
- Representation of d_1 : first column of V_k^T ($k \times n \rightarrow 2 \times 1$) scaled first by S_k :
 $r_d = [0.3, 0.02]^T$
- Query representation: vector average:
 $r_q = [0.13, -0.13]^T / 2 + [0.58, -0.01]^T / 2 = [0.355, -0.07]^T$
- Query-document match: cosine similarity: $\frac{r_q \cdot r_d}{|r_q| \cdot |r_d|} = \frac{0.01205}{0.10879} \approx 0.11$

└ LSA Example

LSA Example

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	0	0	0
...				

- Choose $k = 2$
- Representation of Goethe: fourth row of U_k ($m \times k \rightarrow 1 \times 2$) scaled by S_k : $[0.13, -0.13]^T$
- Representation of devil: fifth row of U_k ($m \times k \rightarrow 1 \times 2$) scaled by S_k : $[0.58, -0.01]^T$
- Representation of d_3 : first column of V_k^T ($k \times n \rightarrow 2 \times 1$) scaled first by S_k : $r_d = [0.3, 0.02]^T$
- Query representation: vector average:
 $r_q = [0.13, -0.13]^T / 2 + [0.58, -0.01]^T / 2 = [0.355, -0.07]^T$
- Query-document match: cosine similarity: $\frac{r_q \cdot r_d}{\|r_q\| \|r_d\|} = \frac{0.01036}{0.01036} \approx 0.11$

- Whether that's a good match or not depends on the ranking and/or threshold

LSA Graphics




Figure 1: Term-document matrix, no ordering, $k = 5$; Source [6]

LSA Graphics




Figure 2: Term-document matrix, group documents, $k = 5$; Source [6]

LSA Graphics




Figure 3: Term-document matrix, group documents+terms, $k = 5$; Source [6]

LSA Code

```
from sklearn.decomposition import TruncatedSVD
from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer(stop_words='english',
                             max_features= 1000,
                             max_df = 0.5,
                             smooth_idf=True)
X = vectorizer.fit_transform(documents)

svd_model = TruncatedSVD(n_components=20)
svd_model.fit(X)
```

Compression: $m \times n \rightarrow m \times k + n \times k + k \times k$

└ LSA Code

LSA Code

```
from sklearn.decomposition import TruncatedSVD
from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer(stop_words='english',
                             max_features= 1000,
                             max_df = 0.5,
                             smooth_idf=True)
X = vectorizer.fit_transform(documents)

svd_model = TruncatedSVD(n_components=20)
svd_model.fit(X)
```

Compression: $m \times n \rightarrow m \times k + n \times k + k \times k$

- max_features takes to top 1000 terms, max_df removes all words which appear in at least half the documents.
- smooth_idf adds one to ever seen term
- The reason it's called Truncated SVD is because it can be used for matrix compression. Instead of transmitting $m \times n$ matrix, we can just transmit the three separate matrices.

Fast SVD

- Naive approach $\det(A - \lambda I) = 0$ solving n -th order polynomial (variable λ)
Eigenvalue Decomposition (EVD), get eigenvectors
- Jacobi rotation [4, 5], Jacobi eigenvalue algorithm [7]:
Create almost a diagonal matrix (bidiagonal): $A = UBV$, $O(mn^2)$
Compute SVD of 2×2 matrixes $O(n^2)$
- Can be parallelized (ARPACK)

Latent Semantic Analysis

- Also called LSI (Latent Semantic Indexing)
- tf-idf is just a weighting scheme (tf, counts)

└ Notes

- tf-idf is not a vital part of LSA, though works well

TODO

- Can be parallelized at the cost of a slightly less accurate approximation

Fast SVD

- Naive approach $\det(A - \lambda I) = 0$ solving n -th order polynomial (variable λ)
- Eigenvector Decomposition (EVD), get eigenvectors
- Jacobi rotation [4, 5], Jacobi eigenvalue algorithm [7]
- Create almost a diagonal matrix (bidiagonal): $A = UBV^T$, $O(mn^2)$
- Compute SVD of 2×2 matrices $O(n^3)$
- Can be parallelized (ARPACK)

Latent Semantic Analysis

- Also called LSI (Latent Semantic Indexing)
- tf-idf is just a weighting scheme (tf, counts)

Considerations

Pros:

- Easy to implement
- Explainable terms
- Quite fast runtime

Cons:

- Only surface dependencies
- SVD is not updatable

Dense Vector Representation

TODO

Resources

- ① Python code:
<https://medium.com/acing-ai/what-is-latent-semantic-analysis-lsa-4d3e2d18417a>
- ② Comprehensive tutorial for LSA+SVD: <https://www.engr.uvic.ca/~seng474/svd.pdf>
- ③ SVD example:
http://web.mit.edu/be.400/www/SVD/Singular_Value_Decomposition.htm
- ④ Computation:
https://en.wikipedia.org/wiki/Singular_value_decomposition#Calculating_the_SVD
- ⑤ Computation: https://www.cs.utexas.edu/users/inderjit/public_papers/HLA_SVD.pdf
- ⑥ Visualization: https://topicmodels.west.uni-koblenz.de/ckling/tmt/svd_ap.html
- ⑦ Computation: https://en.wikipedia.org/wiki/Jacobi_eigenvalue_algorithm
- ⑧ Python code: <https://www.analyticsvidhya.com/blog/2018/10/stepwise-guide-topic-modeling-latent-semantic-analysis/>
- ⑨ Jelinek-Mercer: <http://ctp.di.fct.unl.pt/~jmag/ir/slides/a05%20Language%20models.pdf>
- ⑩ LSI: <https://nlp.stanford.edu/IR-book/html/htmledition/latent-semantic-indexing-1.html>