Information Retrieval + Q&A
(SNLP Tutorial 12)

Vilém Zouhar, Awantee Deshpande, Julius Steuer

13th July, 15th July

Evaluation metrics

- Documents D, queries Q
- System: $Q \to \mathcal{P}(D)$
- For $q \in Q$: retrieved (output), relevant (gold)

- Recall $\frac{|\text{retrieved} \cap \text{relevant}|}{|\text{relevant}|}$
- Precision $\frac{|\text{retrieved} \cap \text{relevant}|}{|\text{retrieved}|}$

Questions?

- When will precision be high? 🍎
- When will recall be high? 🍍

- {Precision,Recall}@k : Retrieve k documents (top k scoring)
- Recall@k $\frac{|\text{retrieved@}k \cap \text{relevant}|}{|\text{relevant}|}$
- Precision@k $\frac{|\text{retrieved@}k \cap \text{relevant}|}{k}$

Evaluation metrics

- Average precision: $AveP(q) = \frac{\sum_k P@k \times rel(k)}{|\text{relevant}|}$
- $rel(k) = \begin{cases} 1 & k\text{-th document relevant} \\ 0 & otherwise \end{cases}$
- Mean average precision $MAP(Q) = \frac{\sum_{q \in Q} AveP(q)}{|Q|}$
-
  - Q can be a "test set"
- F-score $2 \cdot \frac{P \cdot R}{P+R}$

Evaluation metrics

Taking the rank into consideration

- Mean Reciprocal Rank
- $MRR(Q) = \frac{1}{|Q|} \sum_{q \in Q} \frac{1}{\text{rank}_q}$
  rank_q = position of the first relevant document

document	rank	relevant
a	4	+
b	1	
c		
d		+
e	2	+
f	3	

- $Q = \{\text{example}\}, MRR(Q) = \frac{1}{\text{rank}_{example}} = \frac{1}{2}$

Document Retrieval - example

- Query: **Goethe, devil**
- Document:
  A: Wolfgang's idea of the demon Mephistopheles who makes a bet with God
  B: Faust is Wolfgang **Goethe**'s play in German about a pact with the **devil**
  C: **Devil**ishly good lasagne
  D: The impact of **Goethe**'s demon play on the German literature
- How to rank them?
  B (contains the two key words)
  D (Goethe)
  A (Wolfgang - Goethe, Mephistopheles - devil)
  C (unrelated context)
- Can these inferences be made automatically?

Document Retrieval - Bag of Words

- Text must be represented as a vector of numbers
- BoW model requires: i) Vocabulary, ii) Measure of presence of words
- e.g. Vocabulary = {'to', 'be', 'or', 'not', 'question'}
  Document: *to be or not to be*
  BoW representation: {to:2, be:2, or:1, not:1} $\rightarrow$ [1 1 1 1 0]
- Can also store counts
- Disregard grammar, word order

Solution 1 (counts)

- Solution: vector with counts of words:
  `(<the>, <a>, <dog>, <president>, ...)`
  `(57, 68, 0, 2, ...)`
- Issue: representation vectors are enormous
- Issue: longer documents have naturally higher counts
- Issue: useless stop words

Solution 2 (tf)

- Solution: vector with counts of non-stop words, normalized by total words:
  `(<dog>, <president>, <princess>, <thing>, ...)`
  `(0, 0.0003, 0.00001, 0.08, ...)`.
- Issue: some words naturally occur with higher frequency but don't contribute to document meaning (`<thing>`)
- Issue: how do we know which words are useful?

Term Frequency - Inverse Document Frequency

TF-IDF

$$tf(term, doc) = \frac{count_{doc}(term)}{|doc|}$$

$$idf'(term) = \frac{|D|}{df(term)}, idf(term) = \log_2\left(\frac{|D|}{df(term)}\right)$$

$$tf - idf(term, doc) = tf(term, doc) \times idf(term)$$

Augmented TF

$$tf'(term, doc) = 0.5 + 0.5 \cdot \frac{count_{doc}(term)}{max_{term'}\{count_{doc}(term')\}}$$

Term Frequency - Inverse Document Frequency

TF-IDF

$$tf(term, doc) = \frac{count_{doc}(term)}{|doc|}$$

$$idf'(term) = \frac{|D|}{df(term)}, idf(term) = \log_2\left(\frac{|D|}{df(term)}\right)$$

$$tf - idf(term, doc) = tf(term, doc) \times idf(term)$$

Augmented TF

$$tf'(term, doc) = 0.5 + 0.5 \cdot \frac{count_{doc}(term)}{max_{term'}\{count_{doc}(term')\}}$$

- Probability that i-th term occurs k times in the document: $p_{\lambda_i}(k) = e^{-\lambda_i}\frac{\lambda_i^k}{k!}$ (λ_i parameter of the distribution)
- Expected value of occurence: $N \cdot E_i(k) = N \cdot \lambda_i =$ collection frequency$_i$
- Term present at least once: $N \cdot (1 - P_{\lambda_i}(0)) =$ document frequency$_i$

Solution 3 (tf-idf)

- Solution: vector of tf-idf
- Ranking: Cosine similarity between query and document vectors
- Good metrics to determine the significance of a term in a document collection

- Issue: still enormous vectors
- Issue: `demon - Mephistopheles` are equally separate concepts as `devil - lasagne`
- Issue: independent terms assumption

Document Retrieval - Probabilistic Retrieval

- Goal: Find P[R|d,q]
- Ranking: Proportional to relevance odds

$$O(R|d,q) = \frac{P[R|d,q]}{P[\bar{R}|d,q]}$$

- Different probabilistic models calculate these probabilities differently
  e.g. Binary Independence model, Poisson model, BM25

  For Poisson, $P[d|\lambda] = \prod_{t \in V} \frac{e^{-\lambda_t} \cdot \lambda_t^{d_t}}{d_t!}$

Document Retrieval - Statistical Language Model

- Pretend the query was generated by a LM based on the document
- Ranking: Proportional to query likelihood
- $argmax_d \ p(d|q) = argmax_d \ \frac{p(q|d) \cdot p(d)}{p(q)} = argmax_d \ p(q|d) \cdot p(d)$
  $\approx argmax_d \ p_{LM}(q|d) \cdot p(d)$
  $p(d) \approx \frac{1}{|D|}$ or p(d) is query independent
  $\approx argmax_d \ p_{LM}(q|d)$
- Unigram: $p(d|q) \approx \prod_i p_{LM}(q_i|d)$

- LMs can be smoothed, as you remember e.g. Interpolation, Dirichlet smoothing
- Jelinek-Mercer smoothing: $p(q_i|d, D) = \lambda \cdot p(q_i|d) + (1 - \lambda) \cdot p(q_i|D)$
  High λ: documents with all query words (conjunctive)
  Low λ: suitable for long queries (disjunctive)

- Issue: Without word embeddings, no word relatedness
  Query: **Goethe, devil**
  A: Wolfgang's idea of the demon Mephistopheles who makes a bet with God
- Can we model word co-occurence for a topic?

Document Retrieval - Statistical Language Model

- Pretend the query was generated by a LM based on the document
- Ranking: Proportional to query likelihood
- $argmax_d\ p(d|q) = argmax_d\ \frac{p(q|d) \cdot p(d)}{p(q)} = argmax_d\ p(q|d) \cdot p(d)$
  $\approx argmax_d\ p_{LM}(q|d) \cdot p(d)$
  $p(d) \approx \frac{1}{|D|}$ or p(d) is query independent
  $\approx argmax_d\ p_{LM}(q|d)$
- Unigram: $p(d|q) \approx \prod_i p_{LM}(q_i|d)$

- LMs can be smoothed, as you remember e.g. Interpolation, Dirichlet smoothing
- Jelinek-Mercer smoothing: $p(q_i|d, D) = \lambda \cdot p(q_i|d) + (1 - \lambda) \cdot p(q_i|D)$
  High λ: documents with all query words (conjunctive)
  Low λ: suitable for long queries (disjunctive)

- Issue: Without word embeddings, no word relatedness
  Query: **Goethe, devil**
  A: Wolfgang's idea of the demon Mephistopheles who makes a bet with God
- Can we model word co-occurence for a topic?

- Other smoothing schemas exist, like discounting, adding epsilon or linear interpolation between multiple LMs, including zerogram
- Other improvements, such as special grammar, prior knowledge of the document (length), list of synonyms, etc

Solution 4 (Latent Semantic Analysis)

- Assumption: Documents are composed of k latent topics.
- Solution: Perform dimensionality reduction $\rightarrow$ eigenvalues, singular value decomposition
- $A_{i,j} = \#$occurences of term t_i in document d_j

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	0	0	0
Faust	0	1	0	0
Goethe	0	1	0	1
devil	0	1	1	0
demon	1	0	0	1
lasagne	0	0	1	0
German	0	1	0	1

Solution 4 (Latent Semantic Analysis)

- Assumption: Documents are composed of k latent topics.
- Solution: Perform dimensionality reduction $\rightarrow$ eigenvalues, singular value decomposition
- $A_{i,j} = \#$occurences of term t_i in document d_j

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	0	0	0
Faust	0	1	0	0
Goethe	0	1	0	1
devil	0	1	1	0
demon	1	0	0	1
lasagne	0	0	1	0
German	0	1	0	1

The example uses counts, but for better representation of term importance in the document, one would use tf-idf.

Approximation of A

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	**1**	0	**1**
Faust	**1**	1	0	0
Goethe	**1**	1	0	**0**
devil	**1**	1	**0**	**1**
demon	1	**1**	0	1
lasagne	0	0	1	0
German	**1**	1	0	**0**

	c_1	c_2	c_3
Wolfgang	1	0	0
Mephistopheles	0	1	0
Faust	1	0	0
Goethe	1	0	0
devil	0	1	0
demon	0	1	0
lasagne	0	0	1
German	1	0	0

3 latent concepts:
{Goethe (Wolfgang, Faust, German), devil (Mephistopheles, demon), lasagne}

$$d_1 = 1 \times c_1 + 1 \times c_2$$

Approximation of A

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	**1**	0	**1**
Faust	**1**	1	0	0
Goethe	**1**	1	0	**0**
devil	**1**	1	**0**	**1**
demon	1	**1**	0	1
lasagne	0	0	1	0
German	**1**	1	0	**0**

	c_1	c_2	c_3
Wolfgang	1	0	0
Mephistopheles	0	1	0
Faust	1	0	0
Goethe	1	0	0
devil	0	1	0
demon	0	1	0
lasagne	0	0	1
German	1	0	0

3 latent concepts:
{Goethe (Wolfgang, Faust, German), devil (Mephistopheles, demon), lasagne}

$$d_1 = 1 \times c_1 + 1 \times c_2$$

Given k concepts, we try to find such a matrix A', that's as close to the original one, but with every document being a combination of k independent vectors.

Singular Value Decomposition

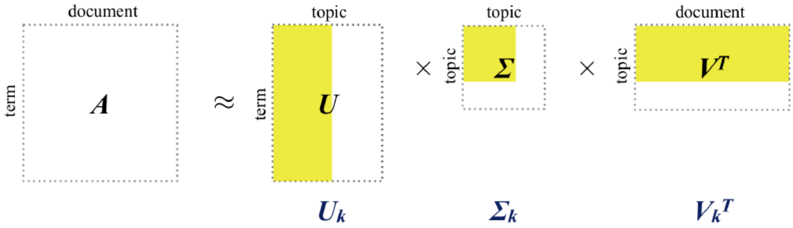

Figure 1: SVD for LSA

- $U =$ eigenvectors of $A^T A$ (# intersection of documents d_i and d_j)
- $V =$ eigenvectors of AA^T (# documents in which both terms t_i and t_j occur)
- $S =$ roots of corresponding eigenvalues of $A^T A$
- $A = USV^T$

LSA

1. Order eigenvalues by descending values ($S_{i,i} > S_{i+1,i+1} \geq 0$) (proof next slide)
2. Take top-k eigenvectors + values (or all above threshold)
3. $A_K = U_K S_K V_K^T$ $[(m \times n), (n \times n), (n \times n)] \rightarrow [(m \times k), (k \times k), (k \times n)]$

- Term $\rightarrow$ latent representation: $U_k S_k$
- Document $\rightarrow$ latent representation: $(S_k V_k^T)^T = V_k S_k^T = V_k S_k$

LSA

1. Order eigenvalues by descending values ($S_{i,i} > S_{i+1,i+1} \geq 0$) (proof next slide)
2. Take top-k eigenvectors + values (or all above threshold)
3. $A_K = U_K S_K V_K^T$ $[(m \times n), (n \times n), (n \times n)] \rightarrow [(m \times k), (k \times k), (k \times n)]$

- Term → latent representation: $U_k S_k$
- Document → latent representation: $(S_k V_k^T)^T = V_k S_k^T = V_k S_k$

- We are free to permute the eigenvalues, so we can order them (together with the vectors) and also we know that the eigenvalues are non-negative
- Therefore we can just take the top-k eigenvalues and replace the rest with zero.
- Essentially this crops the neighbouring matricies to first k columns and first k rows of V^T.

LSA Example

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	0	0	0
...				

- Choose $k = 2$
- Representation of `Goethe`: fourth row of U_k $(m \times k \to 1 \times 2)$ scaled by S_k: $[0.13, -0.13]^T$
- Representation of `devil`: fifth row of U_k $(m \times k \to 1 \times 2)$ scaled by S_k: $[0.58, -0.01]^T$
- Representation of d_1: first column of V_k^T $(k \times n \to 2 \times 1)$ scaled first by S_k: $r_d = [0.3, 0.02]^T$
- Map query to our topic space: $q \to U_k^t \cdot q = q' =$ [0.355, -0.07]^T
- Query-document match: dot product, cosine similarity: $\frac{r_q \cdot r_d}{|r_q| \cdot |r_d|} = \frac{0.01205}{0.10879} \approx 0.11$

LSA Example

	d_1	d_2	d_3	d_4
Wolfgang	1	1	0	0
Mephistopheles	1	0	0	0
...				

- Choose $k = 2$
- Representation of `Goethe`: fourth row of U_k ($m \times k \rightarrow 1 \times 2$) scaled by S_k: $[0.13, -0.13]^T$
- Representation of `devil`: fifth row of U_k ($m \times k \rightarrow 1 \times 2$) scaled by S_k: $[0.58, -0.01]^T$
- Representation of d_1: first column of V_k^T ($k \times n \rightarrow 2 \times 1$) scaled first by S_k: $r_d = [0.3, 0.02]^T$
- Map query to our topic space: $q \rightarrow U_k^t \cdot q = q' =$ [0.355, -0.07]^T
- Query-document match: dot product, cosine similarity: $\frac{r_q \cdot r_d}{|r_q| \cdot |r_d|} = \frac{0.01205}{0.10879} \approx 0.11$

- Whether that's a good match or not depends on the ranking and/or threshold

LSA Graphics

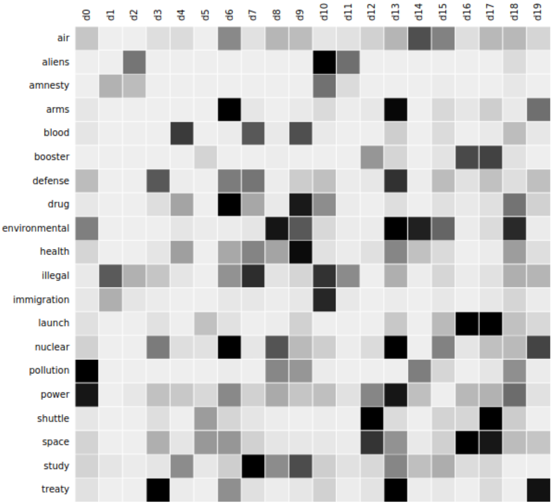

Figure 2: Term-document matrix, no ordering, $k = 5$; Source [6]

LSA Graphics

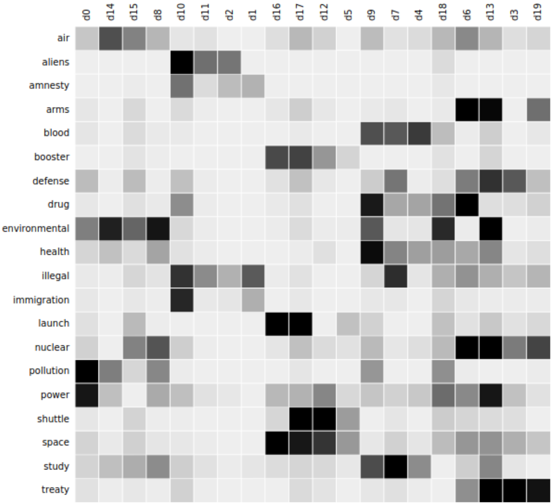

Figure 3: Term-document matrix, group documents, $k = 5$; Source [6]

LSA Graphics

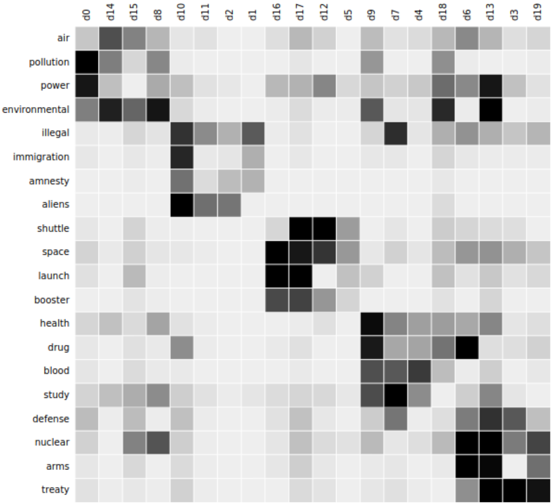

Figure 4: Term-document matrix, group documents+terms, $k = 5$; Source [6]

LSA Code

```
from sklearn.decomposition import TruncatedSVD
from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer(stop_words='english',
    max_features= 1000,
    max_df = 0.5,
    smooth_idf=True)
X = vectorizer.fit_transform(documents)

svd_model = TruncatedSVD(n_components=20)
svd_model.fit(X)
```

Compression: $m \times n \rightarrow m \times k + n \times k + k \times k$

LSA Code

```
from sklearn.decomposition import TruncatedSVD
from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer(stop_words='english',
    max_features= 1000,
    max_df = 0.5,
    smooth_idf=True)
X = vectorizer.fit_transform(documents)

svd_model = TruncatedSVD(n_components=20)
svd_model.fit(X)
```

Compression: $m \times n \to m \times k + n \times k + k \times k$

- max_features takes to top 1000 terms, max_df removes all words which appear in at least half the documents.
- smooth_idf adds one to ever seen term
- The reason it's called Truncated SVD is because it can be used for matrix compression. Instead of transmitting $m \times n$ matrix, we can just transmit the three separate matricies.

Considerations

Pros:

- Easy to implement
- Explainable terms
- Quite fast runtime
- Handles synonymy of words

Cons:

- Only surface dependencies
- Determination of k
- SVD difficult to update

Resources

1. Python code: https://medium.com/acing-ai/what-is-latent-semantic-analysis-lsa-4d3e2d18417a
2. Comprehensive tutorial for LSA+SVD: https://www.engr.uvic.ca/~seng474/svd.pdf
3. SVD example: http://web.mit.edu/be.400/www/SVD/Singular_Value_Decomposition.htm
4. Computation: https://en.wikipedia.org/wiki/Singular_value_decomposition#Calculating_the_SVD
5. Computation: https://www.cs.utexas.edu/users/inderjit/public_papers/HLA_SVD.pdf
6. Visualization: https://topicmodels.west.uni-koblenz.de/ckling/tmt/svd_ap.html
7. Computation: https://en.wikipedia.org/wiki/Jacobi_eigenvalue_algorithm
8. Python code: https://www.analyticsvidhya.com/blog/2018/10/stepwise-guide-topic-modeling-latent-semantic-analysis/
9. Jelinek-Mercer: http://ctp.di.fct.unl.pt/~jmag/ir/slides/a05%20Language%20models.pdf
10. LSI: https://nlp.stanford.edu/IR-book/html/htmledition/latent-semantic-indexing-1.html