

# Statistical Machine Translation

{Phrase-Based, Vanilla Neural}  
(SNLP tutorial)

Vilém Zouhar

March 12, 2021

# Overview

- Task, metrics
- PBMT
  - - Alignment - Phrase extraction
  - - Decoding
    - - - Proof of NP-hardness
    - - - Log-linear model
  - - Alignment - IBM1,2,3,4,5
- NMT
  - - Encoder-Decoder
  - - Embedding
  - - Example
- Homework

# Task

- Given source  $s$ , output target  $t$ :  $\operatorname{argmax}_t \{p(t|s)\}$
- $= \operatorname{argmax}_t \{p(s|t)/p(s) \cdot p(t)\}$
- $= \operatorname{argmax}_t \{p(s|t) \cdot p(t)\}$
- Modelling  $p(s|t)$  is as hard/easy as  $p(t|s)$   
Modelling  $p(t)$  is easier

# Approaches

- RBMT (rule-based)
- EBMT (example-based)
- SMT (statistical)
  - ▶ PBMT (phrase-based)
  - ▶ NMT (neural)




Figure 1: Vauquois Triangle; Source [7]

## └ Approaches

### Approaches

- RBMT (rule-based)
- EBMT (example-based)
- SMT (statistical)
  - PBMT (phrase-based)
  - NMT (neural)



Figure 1: Vasquez Triangle, Source [7]

- Rule-based methods are now a history, though one interesting concept is this triangle.
- Machine translation is split into analysis, transfer and synthesis.
- Example based machine translation is heavily used even today by translation companies, which keep and build translation memories.
- SMT has two parts: phrase based and neural.
- Rule-based are still used in post editing after NMT

# Metrics - BLEU

## BLEU

Reference  $r$ , Candidate  $c$

$$\min(1, \exp(1 - |r|/|c|)) \times (p_1 \cdot p_2 \cdot p_3 \cdot p_4)^{1/4}$$

- Brevity penalty  $\times$  geom. average of precisions
- Score 0-100% (usually without the percentage)

Reference:                   What is the purpose of all this ?

Hypothesis:                What is the meaning of all this ?

matching 3-grams:       (What, is, the), (of, all, this), (all, this, ?)

total 3-grams:            6

→

$$p_3 = 3/6 = 1/2$$

$$p_1 = 7/8, p_2 = 5/7, p_4 = 1/5, \text{BP} = 1, \text{BLEU} = 50.$$

# Metrics

- Many more: ChrF, TER, METEOR
- Every MT metrics faces criticism (~90% correlation with humans [5])
- Human judgement **inconsistent**
- SoTA German-English: ~40
- Strongly depends on:
  - ▶ Tokenization scheme (+10 BLEU(!))
  - ▶ Used corpus
- *Always* use existing BLEU implementation and standardized test sets [6]

# Components of PBMT

- Alignment
  - ▶ Extracting phrases
- Decoding
  - ▶ Covering the source sentence with extracted phrases
  - ▶ Scoring “coverings” using a language model



# Phrase Extraction

- Extract all consistent phrases
- *if you were there*
- *if you were there you would know it now*
- *you would*
- *know it*
- *you would know it now*
- ...
- Extracted phrases have to be “full” - no gaps that are aligned outside of the extraction
- Similar concept to projectivity

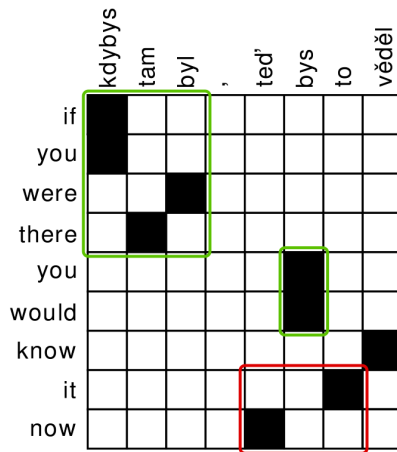


Figure 2: Consistent and inconsistent phrases;  
Source [2]

# Decoding

- Cover the source sentence with extracted phrases

Maria	no	daba	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>		<u>a slap</u>		<u>by</u>		<u>green witch</u>	
	<u>no</u>		<u>slap</u>		<u>to the</u>			
	<u>did not give</u>				<u>to</u>			
					<u>the</u>			
			<u>slap</u>			<u>the witch</u>		

Figure 3: Covering of source sentence; Source [1]

## Beam search

- Start with 0 coverage and keep track of already covered words
- Estimate the cost of the existing phrases (language model) (+ future cost)




Figure 4: Partially expanded beam search; Source [1]

# Beam search - NP-hard

- ① Consider travelling salesman / hamilton circuit
  - ②  $LM(x, y) = -\log dist(x, y)$
  - ③  $LM$ : prohibit repetitions  
 $LM$ : add the distance between the first and the last word
  - ④ Source sentence: NULL-NULL-NULL-NULL-...  
NULL can be covered by any city/node
  - ⑤ Beam search finds the most probable / cheapest ordering:  
NewYork-Boston-Trenton-...
- MT beam search solves the traveling salesman problem  $\rightarrow$  vanilla beam search is NP-hard.
  - Future cost estimation is used + top N hypothesis paths considered (rest pruned).
  - Polynomial, but no optimal solution guarantee.

# PBMT Log-linear Model

- Not statistically sound, but:
- $t = \operatorname{argmax}_t \{p(s|t) \cdot p(t)^2\}$  for more fluent output
- $t = \operatorname{argmax}_t \{p(t|s) \cdot p(t)^2\}$  works equally well

Log-linear model:

- $t = \operatorname{argmax}_t \exp(\sum_{\text{feature } f} \lambda_f f(e, t))$
- Adequacy:  $f_{TM}(e, t) = \log p(t|s)$
- Language model:  $f_{LM}(e, t) = \log p(t)$
- $f_{Phr}(e, t) = \text{number of covering phrases}$ ,  $\lambda_{Phr} = -1$  (e.g.)  
Perhaps we want larger phrases to cover the source sentence

## └ PBMT Log-linear Model

- Not statistically sound, but:
- $t = \operatorname{argmax}_t \{p(s|t) \cdot p(t)^2\}$  for more fluent output
- $t = \operatorname{argmax}_t \{p(t|s) \cdot p(t)^2\}$  works equally well

Log-linear model:

- $t = \operatorname{argmax}_t \exp(\sum_{i=1}^n \lambda_i f_i(s, t))$
- Adequacy:  $f_{TM}(s, t) = \log p(t|s)$
- Language model:  $f_{LM}(s, t) = \log p(t)$
- $f_{ph}(s, t) = \text{number of covering phrases}$ ;  $\lambda_{ph} = -1$  (e.g.)  
Perhaps we want larger phrases to cover the source sentence

- The reason for this notation is that it makes it easy to combine language and adequacy modelling as well as other restrictions.

# Alignment

- Soft alignment: values in the interval  $[0, 1]$  instead of hard decisions  $\{0, 1\}$
- Given an alignment (A), can we construct word translation probabilities (T)?
- Yes:  $T(x|y) = \frac{\sum_{\text{sents } s,t} A_{s,t}(x|y)}{T(\cdot|y)}$
- Given word translation probabilities (T), can we construct word alignment (A)?
- Yes:  $A_{s,t}(x|y) = \frac{T(x|y)}{\sum_{u \in s} T(u|y)}$
- If we start from  $A^1$ , then compute  $T^1$  and then  $A^2$ , will  $A^1 = A^2$ ?
- No, in most cases.
- Main idea behind Expectation-Maximization: change “views” e.g. 5 times.
- Start with  $A_{s,t}^0(x|y) = \frac{1}{|s|}$  (uniform distribution)

## IBM Model 1 Code

```
# expectation
words_prob = np.zeros((len(words2), len(words1)))
for (sent_src, sent_tgt), probs in zip(sents, alignment_probs):
    for word_tgt, propline in zip(sent_tgt, probs):
        for word_src, prob in zip(sent_src, propline):
            words_prob[word_tgt][word_src] += prob

# normalize rows
words_prob = (words_prob.T / np.sum(words_prob, axis=1)).T

# maximization
for sent_i, (sent_src, sent_tgt) in enumerate(sents):
    for pos_src, word_src in enumerate(sent_src):
        for pos_tgt, word_tgt in enumerate(sent_tgt):
            probs[pos_tgt][pos_src] = words_prob[word_tgt][word_src]
# normalize sentence columns
alignment_probs[sent_i] = probs / np.sum(probs, axis=0)
```



# IBM Model 1 - Hard alignment

- At the end, take  $H_{s,t}(y) = \operatorname{argmax}_x (A_{s,t}(x|y))$
- Assumption: every target token is aligned to exactly one source token
- What about alignment between a Slavic language without articles and a Germanic one, with articles?
  - ▶ Czech-German [3]: 24% target tokens unaligned
  - ▶ Czech-German [3]: 1.1 aligned tokens per one target token (excluding unaligned)

Solution:

- Add NULL token to every sentence, then remove alignments to it in post-processing
- Use a different extraction method than argmax (threshold, dynamic threshold, ..)

$$AER = \frac{|An_{sure}| + |An_{poss}|}{|A| + |sure|}$$

## IBM Model 2




Figure 5: Two alignments with equal probability in IBM1

IBM Model 1:  $p(s, \text{align} | t) \propto \prod_{j=1}^{|s|} \text{trans}(s_j | t_{\text{align}(j)})$

IBM Model 2:  $p(s, \text{align} | t) \propto \prod_{j=1}^{|s|} \text{trans}(s_j | t_{\text{align}(j)}) \cdot a(i \rightarrow j, |t|, |s|)$

E.g.  $|\frac{i}{|t|} - \frac{j}{|s|}|$

## └ IBM Model 2




Figure 5: Two alignments with equal probability in IBM1

IBM Model 1:  $p(s, \text{align}(t)) \propto \prod_{i=1}^{|s|} \text{trans}(s_i | t_{\text{align}(i)})$ IBM Model 2:  $p(s, \text{align}(t)) \propto \prod_{i=1}^{|s|} \text{trans}(s_i | t_{\text{align}(i)}) \cdot a(i \rightarrow j, |t|, |s|)$ E.g.  $\frac{1}{|s|} \cdot \frac{1}{|t|}$ 

- The IBM Model 2 introduces a new component, which just scores the probability of just an alignment
- E.g. this way we may force diagonal alignment

## IBM Model 3




Figure 6: Generative process of IBM3; Source [13]

## IBM Model 3




Figure 6: Generative process of IBM3; Source [13]

- The IBM model 3 deals with something called the fertility. That's a concept which captures the fact that for example in Spanish, the English word *slap* is translated to three distinct words *dió una botefada*.
- The fertility of every word can be again estimated using the parallel data.

# IBM Model 4

- Work with classes
- Polish noun-adjective inversion:
- *train station*  $\rightarrow$  *stacja kolejowa*
- $[train:N] [station:N] \rightarrow [stacja:N] [kolejowa:Adj]$  (post-nominal)

## └ IBM Model 4

- Work with classes
- Polish noun-adjective inversion:
  - train station → stacja kolejowa
  - [train:N] [station:N] → [stacja:N] [kolejowa:Adj] (post-nominal)

- IBM Model 4 conditions the alignment probability on word classes of the given word and also the surrounding ones.
- A typical example is the polish noun-adjective inversion, which creates a prior to swap the alignment.

## IBM Model 5

- Alignment context
- Where to align *Übersetzung* if we have low translation priors?




Figure 7: Unfinished alignment process, German  $\rightarrow$  Sotho (South Africa); [14]



## └ IBM Model 5

- Alignment context
- Where to align *Übersetzung* if we have low translation priors?



Figure 7: Unfinished alignment process, German → Sotho (South Africa); [14]

- Finally, IBM model 5 takes the alignment context into consideration.
- That is, it allows placement of words into places that have low translation mass.
- In this example we may have little lexical knowledge
- This can be compensated

# Beyond IBM models

- More heuristics and tricks:
  - ▶ Align everything with levenstein distance at most e.g. 0.1
  - ▶ Align interpunction ( , . ?)
  - ▶ Precision is the biggest issue:  
Compute multiple alignments and output their intersection
- Use existing MT to get translation probabilities
- Transformers (attention scores) for alignment

# NMT - Training

- Traditional SMT pipeline too big + lots of preprocessing
- End-to-end training

① Embed words in one-hot embedding:

dog = (1, 0, 0, 0, ...)

cat = (0, 1, 0, 0, ...)

broccoli = (0, 0, 1, 0, ...)

broccolis = (0, 0, 0, 1, ...)

- ② Feed the whole sentence sequentially into an RNN (vanilla, LSTM, GRU)
- ③ Get a hidden state representing the whole sentence
- ④ The output of this last step is the first translate word (distribution)  
We know the correct word, start accumulating gradient
- ⑤ Push the output word into the hidden state, get next word
- ⑥ Repeat 4.+5. until <EOS> on the training sentence

# NMT - Translation

- Traditional NMT pipeline too big + lots of preprocessing
  - End-to-end training
- 1 Embed words in one-hot embedding
  - 2 Feed the whole sentence sequentially into an RNN (vanilla, LSTM, GRU)
  - 3 Get a hidden state representing the whole sentence
  - 4 The output of this last step is the first translate word (distribution)
  - 5 Push this word into the hidden state, get next word
  - 6 Repeat 4.+5. until <EOS>

In 4. apply beam search or (since we have the probabilities) just take the max.

# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state
- Lots of crucial information is at the beginning of the sentence  
People judge the beginning much more than the rest
- Solution: Feed in the sentence reverse
- Solution: BiRNN
- Solution: Explicit attention
  
- Issue: Whole sentence meaning captured by a single vector
- Solution: Explicit attention mechanism (also alleviates vanishing gradients)  
Transformer architecture (encoder can also be parallelized)
  
- Issue:  $|\text{broccoli-broccolis}|_2^2 = |\text{broccoli-dog}|_2^2$
- Solution: learn word embeddings from monolingual data  
(word2vec: CBOW/skip-gram, Glove)
- Also allows for basic arithmetics: king - man + woman = queen

# Tools

- Alignment:
  - ▶ fast\_align (easy to setup+run, fast, adjusted IBM2) [8]
  - ▶ (M)GIZA++ (more advanced, slightly better results) [9]
- PBMT:
  - ▶ Moses MT [10]
- NMT:
  - ▶ Marian NMT (fast, used by most in WMT, maintained, a bit harder to debug - C++) [11]
  - ▶ Huggingface's transformer (harder to setup, easy Python interop) [12]

# Code

Train:

```
marian \
  --train-sets corpus.en corpus.de \
  --vocabs vocab.en vocab.de \
  --model model.npz
```

Translate:

```
echo "This is a test." | marian-decoder \
  -m model.npz \
  -v vocab.en vocab.de
```

```
> _Das _hi er _ist _ein _Test _.
```



Figure 8: Marian NMT command line options

# Homework

TBD



# References 1

- ❶ PBMT pipeline: <http://www.statmt.org/moses/?n=Moses.Background>
- ❷ Phrase extraction: <https://nlp.fi.muni.cz/en/MachineTranslation>
- ❸ Aligned CSEN corpus: <http://ufal.mff.cuni.cz/czech-english-manual-word-alignment>
- ❹ Capacity of a single \$&!#\* vector: <https://www.aclweb.org/anthology/P18-1198.pdf>
- ❺ BLEU-human annotation correlation:  
<https://www.aclweb.org/anthology/2020.wmt-1.41.pdf>
- ❻ Go-to BLEU implementation: <https://github.com/mjpost/sacreBLEU>
- ❼ SMT course: <http://ufal.mff.cuni.cz/courses/npfl087>
- ❽ fast\_align: [https://github.com/clab/fast\\_align](https://github.com/clab/fast_align)
- ❾ GIZA++: <http://www.statmt.org/moses/giza/GIZA++.html>
- ❿ Moses MT: <http://www.statmt.org/moses/>

## References 2

- ⑪ Marian NMT: <https://marian-nmt.github.io/>
- ⑫ Transformers: <https://huggingface.co/transformers/usage.html>
- ⑬ Stanford lecture:  
<https://web.stanford.edu/class/archive/cs/cs224n/cs224n.1114/handouts/cs224n-lecture-05-2011-MT.pdf>
- ⑭ Alignment visualizer: <https://vilda.net/s/slowalign/>