

# Statistical Machine Translation

{Phrase-Based, Vanilla Neural}  
(SNLP tutorial)

Vilém Zouhar

March 12, 2021

# Overview

- Task, metrics
- PBMT
  - - Alignment - Phrase extraction
  - - Decoding
    - - - Proof of NP-hardness
    - - - Log-linear model
  - - Alignment - IBM1,2,3,4,5
- NMT
  - - Encoder-Decoder
  - - Embedding
  - - Example
- Homework

# Task

- Given source  $s$ , output target  $t$ :  $\operatorname{argmax}_t \{p(t|s)\}$

# Task

- Given source  $s$ , output target  $t$ :  $\operatorname{argmax}_t \{p(t|s)\}$
- $= \operatorname{argmax}_t \{p(s|t)/p(s) \cdot p(t)\}$

# Task

- Given source  $s$ , output target  $t$ :  $\operatorname{argmax}_t \{p(t|s)\}$
- $= \operatorname{argmax}_t \{p(s|t)/p(s) \cdot p(t)\}$
- $= \operatorname{argmax}_t \{p(s|t) \cdot p(t)\}$

# Task

- Given source  $s$ , output target  $t$ :  $\operatorname{argmax}_t \{p(t|s)\}$
- $= \operatorname{argmax}_t \{p(s|t)/p(s) \cdot p(t)\}$
- $= \operatorname{argmax}_t \{p(s|t) \cdot p(t)\}$
- Modelling  $p(s|t)$  is as hard/easy as  $p(t|s)$   
Modelling  $p(t)$  is easier

# Approaches

- RBMT (rule-based)
- EBMT (example-based)
- SMT (statistical)
  - ▶ PBMT (phrase-based)
  - ▶ NMT (neural)

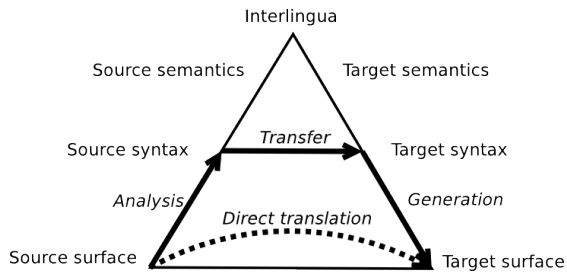


Figure 1: Vauquois Triangle; Source [7]

# Metrics - BLEU

## BLEU

Reference  $r$ , Candidate  $c$

$$\min(1, \exp(1 - |r|/|c|)) \times (p_1 \cdot p_2 \cdot p_3 \cdot p_4)^{1/4}$$

- Brevity penalty  $\times$  geom. average of precisions
- Score 0-100% (usually without the percentage)

Reference:                   What is the purpose of all this ?

Hypothesis:                What is the meaning of all this ?

matching 3-grams:       (What, is, the), (of, all, this), (all, this, ?)

total 3-grams:            6

→

$$p_3 = 3/6 = 1/2$$

$$p_1 = 7/8, p_2 = 5/7, p_4 = 1/5, \text{BP} = 1, \text{BLEU} = 50.$$



# Metrics

- Many more: ChrF, TER, METEOR
- Every MT metrics faces criticism (~90% correlation with humans [5])
- Human judgement **inconsistent**
- SoTA German-English: ~40
- Strongly depends on:
  - ▶ Tokenization scheme (+10 BLEU(!))
  - ▶ Used corpus
- *Always* use existing BLEU implementation and standardized test sets [6]

# Components of PBMT

- Alignment
  - ▶ Extracting phrases
- Decoding
  - ▶ Covering the source sentence with extracted phrases
  - ▶ Scoring “coverings” using a language model

## Phrase Extraction

- Extract all consistent phrases
- *if you were there*
- *if you were there you would know it now*
- *you would*
- *know it*
- *you would know it now*
- ...
- Extracted phrases have to be “full” - no gaps that are aligned outside of the extraction
- Similar concept to projectivity




Figure 2: Consistent and inconsistent phrases;  
Source [2]

# Decoding

- Cover the source sentence with extracted phrases

Maria	no	daba	una	bofetada	a	la	bruja	verde
<u>Mary</u>	<u>not</u>	<u>give</u>	<u>a</u>	<u>slap</u>	<u>to</u>	<u>the</u>	<u>witch</u>	<u>green</u>
	<u>did not</u>		<u>a slap</u>		<u>by</u>		<u>green witch</u>	
	<u>no</u>		<u>slap</u>		<u>to the</u>			
	<u>did not give</u>				<u>to</u>			
					<u>the</u>			
			<u>slap</u>			<u>the witch</u>		

Figure 3: Covering of source sentence; Source [1]

## Beam search

- Start with 0 coverage and keep track of already covered words
- Estimate the cost of the existing phrases (language model) (+ future cost)




Figure 4: Partially expanded beam search; Source [1]

# Beam search - NP-hard

- ① Consider travelling salesman / hamilton circuit
  - ②  $LM(x, y) = -\log dist(x, y)$
  - ③  $LM$ : prohibit repetitions  
 $LM$ : add the distance between the first and the last word
  - ④ Source sentence: NULL-NULL-NULL-NULL-...  
NULL can be covered by any city/node
  - ⑤ Beam search finds the most probable / cheapest ordering:  
NewYork-Boston-Trenton-...
- MT beam search solves the traveling salesman problem  $\rightarrow$  vanilla beam search is NP-hard.
  - Future cost estimation is used + top N hypothesis paths considered (rest pruned).
  - Polynomial, but no optimal solution guarantee.

# PBMT Log-linear Model

- Not statistically sound, but:
- $t = \operatorname{argmax}_t \{p(s|t) \cdot p(t)^2\}$  for more fluent output
- $t = \operatorname{argmax}_t \{p(t|s) \cdot p(t)^2\}$  works equally well

# PBMT Log-linear Model

- Not statistically sound, but:
- $t = \operatorname{argmax}_t \{p(s|t) \cdot p(t)^2\}$  for more fluent output
- $t = \operatorname{argmax}_t \{p(t|s) \cdot p(t)^2\}$  works equally well

Log-linear model:

- $t = \operatorname{argmax}_t \exp(\sum_{\text{feature } f} \lambda_f f(e, t))$



# PBMT Log-linear Model

- Not statistically sound, but:
- $t = \operatorname{argmax}_t \{p(s|t) \cdot p(t)^2\}$  for more fluent output
- $t = \operatorname{argmax}_t \{p(t|s) \cdot p(t)^2\}$  works equally well

Log-linear model:

- $t = \operatorname{argmax}_t \exp(\sum_{\text{feature } f} \lambda_f f(e, t))$
- Adequacy:  $f_{TM}(e, t) = \log p(t|s)$

# PBMT Log-linear Model

- Not statistically sound, but:
- $t = \operatorname{argmax}_t \{p(s|t) \cdot p(t)^2\}$  for more fluent output
- $t = \operatorname{argmax}_t \{p(t|s) \cdot p(t)^2\}$  works equally well

Log-linear model:

- $t = \operatorname{argmax}_t \exp(\sum_{\text{feature } f} \lambda_f f(e, t))$
- Adequacy:  $f_{TM}(e, t) = \log p(t|s)$
- Language model:  $f_{LM}(e, t) = \log p(t)$

# PBMT Log-linear Model

- Not statistically sound, but:
- $t = \operatorname{argmax}_t \{p(s|t) \cdot p(t)^2\}$  for more fluent output
- $t = \operatorname{argmax}_t \{p(t|s) \cdot p(t)^2\}$  works equally well

Log-linear model:

- $t = \operatorname{argmax}_t \exp(\sum_{\text{feature } f} \lambda_f f(e, t))$
- Adequacy:  $f_{TM}(e, t) = \log p(t|s)$
- Language model:  $f_{LM}(e, t) = \log p(t)$
- $f_{Phr}(e, t) = \text{number of covering phrases}, \lambda_{Phr} = -1$  (e.g.)  
Perhaps we want larger phrases to cover the source sentence

# Alignment

- Soft alignment: values in the interval  $[0, 1]$  instead of hard decisions  $\{0, 1\}$

# Alignment

- Soft alignment: values in the interval  $[0, 1]$  instead of hard decisions  $\{0, 1\}$
- Given an alignment (A), can we construct word translation probabilities (T)?

# Alignment

- Soft alignment: values in the interval  $[0, 1]$  instead of hard decisions  $\{0, 1\}$
- Given an alignment (A), can we construct word translation probabilities (T)?
- Yes:  $T(x|y) = \frac{\sum_{\text{sents } s,t} A_{s,t}(x|y)}{T(\cdot|y)}$

# Alignment

- Soft alignment: values in the interval  $[0, 1]$  instead of hard decisions  $\{0, 1\}$
- Given an alignment (A), can we construct word translation probabilities (T)?
- Yes:  $T(x|y) = \frac{\sum_{\text{sents } s,t} A_{s,t}(x|y)}{T(\cdot|y)}$
- Given word translation probabilities (T), can we construct word alignment (A)?

# Alignment

- Soft alignment: values in the interval  $[0, 1]$  instead of hard decisions  $\{0, 1\}$
- Given an alignment (A), can we construct word translation probabilities (T)?
- Yes:  $T(x|y) = \frac{\sum_{\text{sents } s,t} A_{s,t}(x|y)}{T(\cdot|y)}$
- Given word translation probabilities (T), can we construct word alignment (A)?
- Yes:  $A_{s,t}(x|y) = \frac{T(x|y)}{\sum_{u \in s} T(u|y)}$



# Alignment

- Soft alignment: values in the interval  $[0, 1]$  instead of hard decisions  $\{0, 1\}$
- Given an alignment (A), can we construct word translation probabilities (T)?
- Yes:  $T(x|y) = \frac{\sum_{\text{sents } s,t} A_{s,t}(x|y)}{T(\cdot|y)}$
- Given word translation probabilities (T), can we construct word alignment (A)?
- Yes:  $A_{s,t}(x|y) = \frac{T(x|y)}{\sum_{u \in s} T(u|y)}$
- If we start from  $A^1$ , then compute  $T^1$  and then  $A^2$ , will  $A^1 = A^2$ ?

# Alignment

- Soft alignment: values in the interval  $[0, 1]$  instead of hard decisions  $\{0, 1\}$
- Given an alignment (A), can we construct word translation probabilities (T)?
- Yes:  $T(x|y) = \frac{\sum_{\text{sents } s,t} A_{s,t}(x|y)}{T(\cdot|y)}$
- Given word translation probabilities (T), can we construct word alignment (A)?
- Yes:  $A_{s,t}(x|y) = \frac{T(x|y)}{\sum_{u \in s} T(u|y)}$
- If we start from  $A^1$ , then compute  $T^1$  and then  $A^2$ , will  $A^1 = A^2$ ?
- No, in most cases.

# Alignment

- Soft alignment: values in the interval  $[0, 1]$  instead of hard decisions  $\{0, 1\}$
- Given an alignment (A), can we construct word translation probabilities (T)?
- Yes:  $T(x|y) = \frac{\sum_{\text{sents } s,t} A_{s,t}(x|y)}{T(\cdot|y)}$
- Given word translation probabilities (T), can we construct word alignment (A)?
- Yes:  $A_{s,t}(x|y) = \frac{T(x|y)}{\sum_{u \in s} T(u|y)}$
- If we start from  $A^1$ , then compute  $T^1$  and then  $A^2$ , will  $A^1 = A^2$ ?
- No, in most cases.
- Main idea behind Expectation-Maximization: change “views” e.g. 5 times.

# Alignment

- Soft alignment: values in the interval  $[0, 1]$  instead of hard decisions  $\{0, 1\}$
- Given an alignment (A), can we construct word translation probabilities (T)?
- Yes:  $T(x|y) = \frac{\sum_{\text{sents } s,t} A_{s,t}(x|y)}{T(\cdot|y)}$
- Given word translation probabilities (T), can we construct word alignment (A)?
- Yes:  $A_{s,t}(x|y) = \frac{T(x|y)}{\sum_{u \in s} T(u|y)}$
- If we start from  $A^1$ , then compute  $T^1$  and then  $A^2$ , will  $A^1 = A^2$ ?
- No, in most cases.
- Main idea behind Expectation-Maximization: change “views” e.g. 5 times.
- Start with  $A_{s,t}^0(x|y) = \frac{1}{|s|}$  (uniform distribution)

## IBM Model 1 Code

```
# expectation
words_prob = np.zeros((len(words2), len(words1)))
for (sent_src, sent_tgt), probs in zip(sents, alignment_probs):
    for word_tgt, propline in zip(sent_tgt, probs):
        for word_src, prob in zip(sent_src, propline):
            words_prob[word_tgt][word_src] += prob

# normalize rows
words_prob = (words_prob.T / np.sum(words_prob, axis=1)).T

# maximization
for sent_i, (sent_src, sent_tgt) in enumerate(sents):
    for pos_src, word_src in enumerate(sent_src):
        for pos_tgt, word_tgt in enumerate(sent_tgt):
            probs[pos_tgt][pos_src] = words_prob[word_tgt][word_src]
# normalize sentence columns
alignment_probs[sent_i] = probs / np.sum(probs, axis=0)
```

# IBM Model 1 - Hard alignment

- At the end, take  $H_{s,t}(y) = \operatorname{argmax}_x (A_{s,t}(x|y))$
- Assumption: every target token is aligned to exactly one source token
- What about alignment between a Slavic language without articles and a Germanic one, with articles?
  - ▶ Czech-German [3]: 24% target tokens unaligned
  - ▶ Czech-German [3]: 1.1 aligned tokens per one target token (excluding unaligned)

# IBM Model 1 - Hard alignment

- At the end, take  $H_{s,t}(y) = \operatorname{argmax}_x (A_{s,t}(x|y))$
- Assumption: every target token is aligned to exactly one source token
- What about alignment between a Slavic language without articles and a Germanic one, with articles?
  - ▶ Czech-German [3]: 24% target tokens unaligned
  - ▶ Czech-German [3]: 1.1 aligned tokens per one target token (excluding unaligned)

Solution:

- Add NULL token to every sentence, then remove alignments to it in post-processing
- Use a different extraction method than argmax (threshold, dynamic threshold, ..)

# IBM Model 1 - Hard alignment

- At the end, take  $H_{s,t}(y) = \operatorname{argmax}_x (A_{s,t}(x|y))$
- Assumption: every target token is aligned to exactly one source token
- What about alignment between a Slavic language without articles and a Germanic one, with articles?
  - ▶ Czech-German [3]: 24% target tokens unaligned
  - ▶ Czech-German [3]: 1.1 aligned tokens per one target token (excluding unaligned)

Solution:

- Add NULL token to every sentence, then remove alignments to it in post-processing
- Use a different extraction method than argmax (threshold, dynamic threshold, ..)

$$AER = \frac{|An_{sure}| + |An_{poss}|}{|A| + |sure|}$$



## IBM Model 2




Figure 5: Two alignments with equal probability in IBM1

IBM Model 1:  $p(s, \text{align} | t) \propto \prod_{j=1}^{|s|} \text{trans}(s_j | t_{\text{align}(j)})$

IBM Model 2:  $p(s, \text{align} | t) \propto \prod_{j=1}^{|s|} \text{trans}(s_j | t_{\text{align}(j)}) \cdot a(i \rightarrow j, |t|, |s|)$

E.g.  $|\frac{i}{|t|} - \frac{j}{|s|}|$

## IBM Model 3

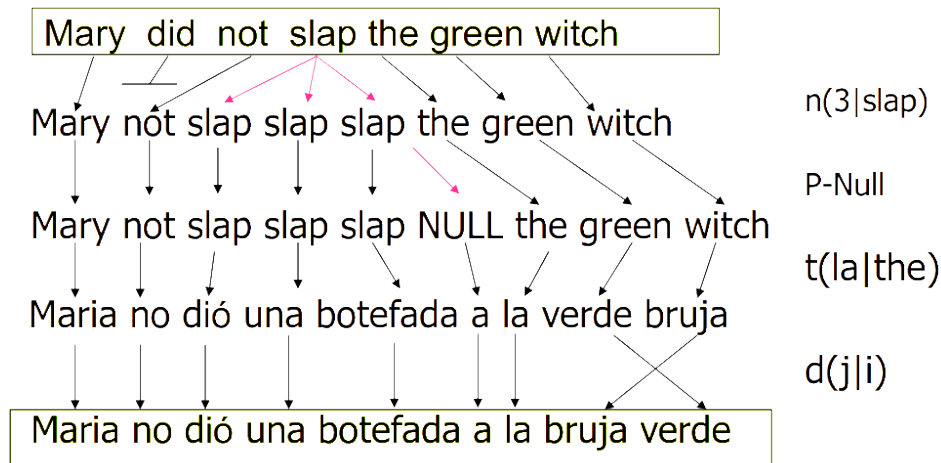


Figure 6: Generative process of IBM3; Source [13]

# IBM Model 4

- Work with classes
- Polish noun-adjective inversion:
- *train station*  $\rightarrow$  *stacja kolejowa*
- $[train:N] [station:N] \rightarrow [stacja:N] [kolejowa:Adj]$  (post-nominal)

## IBM Model 5

- Alignment context
- Where to align *Übersetzung* if we have low translation priors?




Figure 7: Unfinished alignment process, German  $\rightarrow$  Sotho (South Africa); [14]

# Beyond IBM models

- More heuristics and tricks:

# Beyond IBM models

- More heuristics and tricks:

# Beyond IBM models

- More heuristics and tricks:
  - ▶ Align everything with levenstein distance at most e.g. 0.1

# Beyond IBM models

- More heuristics and tricks:
  - ▶ Align everything with levenstein distance at most e.g. 0.1



# Beyond IBM models

- More heuristics and tricks:
  - ▶ Align everything with levenstein distance at most e.g. 0.1
  - ▶ Align interpunction (, . ?)

# Beyond IBM models

- More heuristics and tricks:
  - ▶ Align everything with levenstein distance at most e.g. 0.1
  - ▶ Align interpunction (, . ?)

# Beyond IBM models

- More heuristics and tricks:
  - ▶ Align everything with levenstein distance at most e.g. 0.1
  - ▶ Align interpunction (, . ?)
  - ▶ Precision is the biggest issue:  
Compute multiple alignments and output their intersection

# Beyond IBM models

- More heuristics and tricks:
  - ▶ Align everything with levenstein distance at most e.g. 0.1
  - ▶ Align interpunction ( , . ?)
  - ▶ Precision is the biggest issue:  
Compute multiple alignments and output their intersection
- Use existing MT to get translation probabilities

# Beyond IBM models

- More heuristics and tricks:
  - ▶ Align everything with levenstein distance at most e.g. 0.1
  - ▶ Align interpunction ( , . ?)
  - ▶ Precision is the biggest issue:  
Compute multiple alignments and output their intersection
- Use existing MT to get translation probabilities
- Transformers (attention scores) for alignment

# NMT - Training

- Traditional SMT pipeline too big + lots of preprocessing
- End-to-end training

① Embed words in one-hot embedding:

dog = (1, 0, 0, 0, ...)

cat = (0, 1, 0, 0, ...)

broccoli = (0, 0, 1, 0, ...)

broccolis = (0, 0, 0, 1, ...)

- ② Feed the whole sentence sequentially into an RNN (vanilla, LSTM, GRU)
- ③ Get a hidden state representing the whole sentence
- ④ The output of this last step is the first translate word (distribution)  
We know the correct word, start accumulating gradient
- ⑤ Push the output word into the hidden state, get next word
- ⑥ Repeat 4.+5. until <EOS> on the training sentence

# NMT - Translation

- Traditional NMT pipeline too big + lots of preprocessing
  - End-to-end training
- 1 Embed words in one-hot embedding
  - 2 Feed the whole sentence sequentially into an RNN (vanilla, LSTM, GRU)
  - 3 Get a hidden state representing the whole sentence
  - 4 The output of this last step is the first translate word (distribution)
  - 5 Push this word into the hidden state, get next word
  - 6 Repeat 4.+5. until <EOS>

In 4. apply beam search or (since we have the probabilities) just take the max.

# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state



# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state
- Lots of crucial information is at the beginning of the sentence  
People judge the beginning much more than the rest

# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state
- Lots of crucial information is at the beginning of the sentence  
People judge the beginning much more than the rest
- Solution: Feed in the sentence reverse

# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state
- Lots of crucial information is at the beginning of the sentence  
People judge the beginning much more than the rest
- Solution: Feed in the sentence reverse
- Solution: BiRNN

# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state
- Lots of crucial information is at the beginning of the sentence  
People judge the beginning much more than the rest
- Solution: Feed in the sentence reverse
- Solution: BiRNN
- Solution: Explicit attention

# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state
- Lots of crucial information is at the beginning of the sentence  
People judge the beginning much more than the rest
- Solution: Feed in the sentence reverse
- Solution: BiRNN
- Solution: Explicit attention
- Issue: Whole sentence meaning captured by a single vector

# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state
- Lots of crucial information is at the beginning of the sentence  
People judge the beginning much more than the rest
- Solution: Feed in the sentence reverse
- Solution: BiRNN
- Solution: Explicit attention
  
- Issue: Whole sentence meaning captured by a single vector
- Solution: Explicit attention mechanism (also alleviates vanishing gradients)  
Transformer architecture (encoder can also be parallelized)

# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state
- Lots of crucial information is at the beginning of the sentence  
People judge the beginning much more than the rest
- Solution: Feed in the sentence reverse
- Solution: BiRNN
- Solution: Explicit attention
  
- Issue: Whole sentence meaning captured by a single vector
- Solution: Explicit attention mechanism (also alleviates vanishing gradients)  
Transformer architecture (encoder can also be parallelized)
  
- Issue:  $|\text{broccoli-broccolis}|_2^2 = |\text{broccoli-dog}|_2^2$

# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state
- Lots of crucial information is at the beginning of the sentence  
People judge the beginning much more than the rest
- Solution: Feed in the sentence reverse
- Solution: BiRNN
- Solution: Explicit attention
  
- Issue: Whole sentence meaning captured by a single vector
- Solution: Explicit attention mechanism (also alleviates vanishing gradients)  
Transformer architecture (encoder can also be parallelized)
  
- Issue:  $|\text{broccoli-broccolis}|_2^2 = |\text{broccoli-dog}|_2^2$
- Solution: learn word embeddings from monolingual data  
(word2vec: CBOW/skip-gram, Glove)



# NMT - Issues

- Issue: The end of the sentence is represented best in the single hidden state
- Lots of crucial information is at the beginning of the sentence  
People judge the beginning much more than the rest
- Solution: Feed in the sentence reverse
- Solution: BiRNN
- Solution: Explicit attention
  
- Issue: Whole sentence meaning captured by a single vector
- Solution: Explicit attention mechanism (also alleviates vanishing gradients)  
Transformer architecture (encoder can also be parallelized)
  
- Issue:  $|\text{broccoli-broccolis}|_2^2 = |\text{broccoli-dog}|_2^2$
- Solution: learn word embeddings from monolingual data  
(word2vec: CBOW/skip-gram, Glove)
- Also allows for basic arithmetics: king - man + woman = queen

# Tools

- Alignment:
  - ▶ fast\_align (easy to setup+run, fast, adjusted IBM2) [8]
  - ▶ (M)GIZA++ (more advanced, slightly better results) [9]
- PBMT:
  - ▶ Moses MT [10]
- NMT:
  - ▶ Marian NMT (fast, used by most in WMT, maintained, a bit harder to debug - C++) [11]
  - ▶ Huggingface's transformer (harder to setup, easy Python interop) [12]

## Code

Train:

```
marian \  
  --train-sets corpus.en corpus.de \  
  --vocabs vocab.en vocab.de \  
  --model model.npz
```

## Code

Train:

```
marian \  
  --train-sets corpus.en corpus.de \  
  --vocabs vocab.en vocab.de \  
  --model model.npz
```

Translate:

```
echo "This is a test." | marian-decoder \  
  -m model.npz \  
  -v vocab.en vocab.de
```

```
> _Das _hi er _ist _ein _Test _.
```

# Code

Train:

```
marian \
  --train-sets corpus.en corpus.de \
  --vocabs vocab.en vocab.de \
  --model model.npz
```

Translate:

```
echo "This is a test." | marian-decoder \
  -m model.npz \
  -v vocab.en vocab.de
```

```
> _Das _hi er _ist _ein _Test _.
```



The image shows a screenshot of the command-line help for the Marian NMT software. It is organized into three main sections: 'Global options', 'Training options', and 'Decoding options'. Each section contains a list of flags and their descriptions. For example, under 'Global options', there are flags for verbosity, quiet mode, and help. Under 'Training options', there are flags for training sets, vocabularies, model file, and various training parameters like beam size and sampling. Under 'Decoding options', there are flags for the model file, vocabularies, and decoding parameters like beam size and sampling.

Figure 8: Marian NMT command line options

# Homework

TBD

# References 1

- ❶ PBMT pipeline: <http://www.statmt.org/moses/?n=Moses.Background>
- ❷ Phrase extraction: <https://nlp.fi.muni.cz/en/MachineTranslation>
- ❸ Aligned CSEN corpus: <http://ufal.mff.cuni.cz/czech-english-manual-word-alignment>
- ❹ Capacity of a single \$&!#\* vector: <https://www.aclweb.org/anthology/P18-1198.pdf>
- ❺ BLEU-human annotation correlation:  
<https://www.aclweb.org/anthology/2020.wmt-1.41.pdf>
- ❻ Go-to BLEU implementation: <https://github.com/mjpost/sacreBLEU>
- ❼ SMT course: <http://ufal.mff.cuni.cz/courses/npfl087>
- ❽ fast\_align: [https://github.com/clab/fast\\_align](https://github.com/clab/fast_align)
- ❾ GIZA++: <http://www.statmt.org/moses/giza/GIZA++.html>
- ❿ Moses MT: <http://www.statmt.org/moses/>

## References 2

- ⑪ Marian NMT: <https://marian-nmt.github.io/>
- ⑫ Transformers: <https://huggingface.co/transformers/usage.html>
- ⑬ Stanford lecture:  
<https://web.stanford.edu/class/archive/cs/cs224n/cs224n.1114/handouts/cs224n-lecture-05-2011-MT.pdf>
- ⑭ Alignment visualizer: <https://vilda.net/s/slowalign/>